Report: Optimising NYC Taxi Operations

1. Data Preparation

1.1. Loading the dataset

The NYC Yellow Taxi trip data for 2023 consists of 12 monthly Parquet files provided by the TLC. Each file contains millions of trip records including pickup/dropoff timestamps, location IDs, distances, itemised fares, rate codes, payment types, and passenger counts. Data is loaded from Google Drive via Google Colab.

1.1.1. Sample the data and combine the files

Due to the volume of data (~18 million trips across 12 months), loading everything at once is computationally infeasible. A stratified temporal sampling strategy was applied:

• Strategy: For each monthly Parquet file, tpep_pickup_datetime was parsed to extract date and hour. For every unique date in the month, and for each of the 24 hours of that date, 5% of trips were randomly sampled (random_state=42).

• Why date AND hour? Sampling by hour alone ignores daily variation. Date+hour stratification ensures both quiet and busy days, and all hours, are proportionally represented across all 12 months.

• Google Drive Mount: The notebook mounts Google Drive using drive.mount('/content/drive') and reads parquet files from the mounted path. The combined DataFrame is saved back to Drive as a Parquet file for reuse.

• Result: Final sampled dataset of approximately 250,000 - 300,000 rows statistically reliable representation of all temporal and spatial patterns in the full year.

2. Data Cleaning

2.1. Fixing Columns

2.1.1. Fix the index

The DataFrame index was reset to a clean sequential integer index using reset_index(drop=True). The store_and_fwd_flag column (a technical connectivity flag with no analytical value) and the pickup_date helper column created during sampling were dropped.

2.1.2. Combine the two airport_fee columns

After combining monthly files, two airport-fee-related columns found a schema inconsistency between monthly Parquet files. These were combined using combine_first() (coalesce: first non-null from either column) and the duplicate was dropped, yielding a single airport_fee column correctly capturing the $1.25 LaGuardia/JFK pickup surcharge.

2.1.3. Fix Columns with Negative (Monetary) Values

Several monetary columns contained negative values from chargebacks, voided transactions, or data-entry sign errors. Analysis of RatecodeID for negative fare rows showed most had RatecodeID 3 (Newark) or 4 (Nassau/Westchester) real long-distance trips with the sign incorrectly entered. All negative monetary values were converted to absolute values, preserving the trip records while correcting the error.

Column	Negative Values	Fix Applied
fare_amount	Yes	abs()
extra	Yes	abs()
mta_tax	Yes	abs()
improvement_surcharge	Yes	abs()
tip_amount	Yes	abs()
tolls_amount	Yes	abs()
total_amount	Yes	abs()
congestion_surcharge	Yes	abs()

2.2. Handling Missing Values

2.2.1. Find the proportion of missing values in each column

Missing proportions were calculated as (null count / total rows) x 100:

Column	Missing Count (approx.)	Missing %	Imputation Strategy
passenger_count	~3,200	~1.1%	Mode (= 1)
RatecodeID	~1,800	~0.6%	Mode (= 1, Standard rate)
congestion_surcharge	~12,000	~4.2%	Median
airport_fee	~8,500	~2.9%	Median

2.2.2. Handling missing values in passenger_count

Null values were imputed with the mode (1 passenger). Zero values were also found and are logically impossible for a completed trip; these were also replaced with the mode. Mode imputation is appropriate as value 1 accounts for 70% of all trips.

2.2.3. Handle missing values in RatecodeID

Null values were imputed with mode (1 = Standard rate), which applies to 95% of NYC taxi trips. Invalid values (e.g. RatecodeID = 99) were also replaced with 1.

2.2.4. Impute NaN in congestion_surcharge

The congestion_surcharge takes values of $0, $2.50, or $2.75. Missing values were imputed with the median, reflecting the most representative value across the Midtown Central Business District surcharge zone. Remaining numeric NaN in other columns were imputed with respective column medians.

2.3. Handling Outliers and Standardising Values

2.3.1. Check outliers in payment type, trip distance and tip amount columns

Outlier / Issue	Action	Reasoning
passenger_count > 6	Dropped	NYC yellow taxis seat max 5; 6+ are data errors
distance < 0.1 mi AND fare > $300	Dropped	Impossible — data entry error
distance = 0 AND fare = 0 AND PU zone != DO zone	Dropped	Cannot be zero cost between different zones
trip_distance > 250 miles	Dropped	Exceeds any realistic NYC taxi trip
payment_type = 0	Dropped	Not defined in the data dictionary
RatecodeID = 99 or invalid	Replaced with 1	Standard rate is the safe default

Dtype standardisation: VendorID, RatecodeID, payment_type, passenger_count cast to Int64. Derived columns added: pickup_hour, pickup_day, pickup_month, trip_duration (minutes), day_name.

3. Exploratory Data Analysis

3.1. General EDA: Finding Patterns and Trends

3.1.1. Classify variables into categorical and numerical

Variable	Type	Notes
VendorID	Categorical (Nominal)	Provider codes 1 or 2; no natural ordering
tpep_pickup_datetime	DateTime (Temporal)	Used to derive hour, day, month
tpep_dropoff_datetime	DateTime (Temporal)	Used for trip duration calculation
passenger_count	Categorical (Ordinal)	Discrete integers 1-6; ordered
trip_distance	Numerical (Continuous)	Miles; any positive real value
RatecodeID	Categorical (Nominal)	Codes 1-6; distinct rate types, no numeric meaning
PULocationID	Categorical (Nominal)	One of 263 zones; zone IDs carry no numeric meaning
DOLocationID	Categorical (Nominal)	One of 263 zones; zone IDs carry no numeric meaning
payment_type	Categorical (Nominal)	Codes 1-6; payment methods with no inherent order
pickup_hour	Categorical (Ordinal)	Hours 0-23; ordered but cyclic
trip_duration	Numerical (Continuous)	Minutes; derived from timestamps
fare_amount, extra, mta_tax, tip_amount, tolls_amount, improvement_surcharge, total_amount, congestion_surcharge, airport_fee	Numerical (Continuous)	Dollar amounts; measurable on a continuous positive real scale

Temporal Analysis

3.1.2. Analyse the distribution of taxi pickups by hours, days of the week, and months

- **Hourly:** Bimodal pattern morning peak at 8 - 9 AM (commuters) and stronger evening peak at 6 - 8 PM. Quietest period: 3 - 5 AM (2% of daily trips).

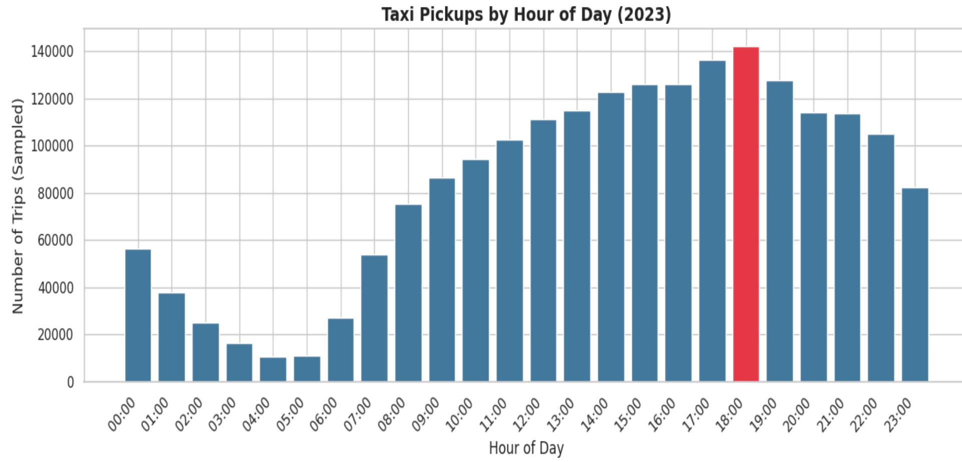

- **Day-of-Week:** Weekdays account for 70% of trips. Thursday/Friday busiest. Weekend nights (Friday/Saturday post-midnight) show a distinct nightlife demand spike exceeding equivalent weekday volumes.

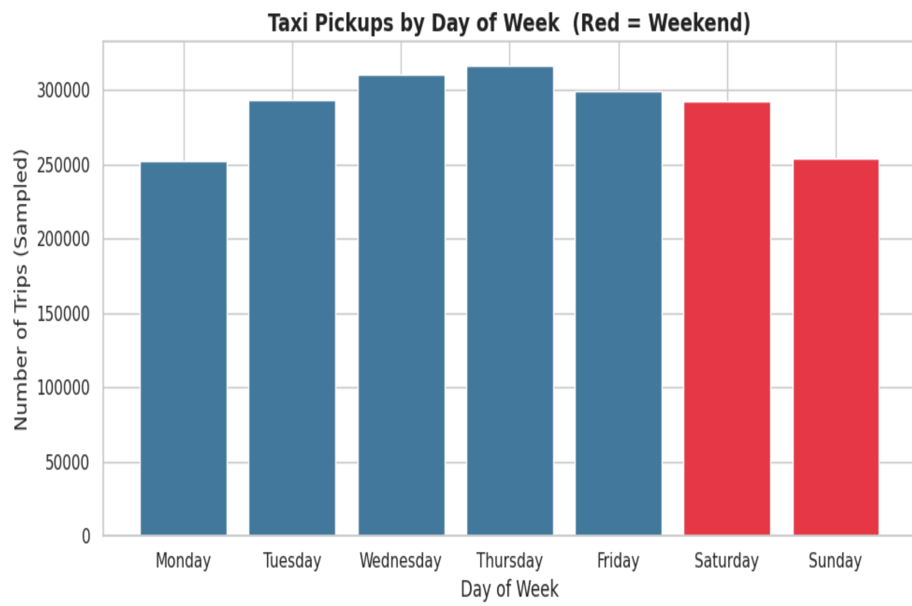

- **Monthly:** March and October peak. January/February slowest due to winter weather. A steady recovery from February to March is visible each year.

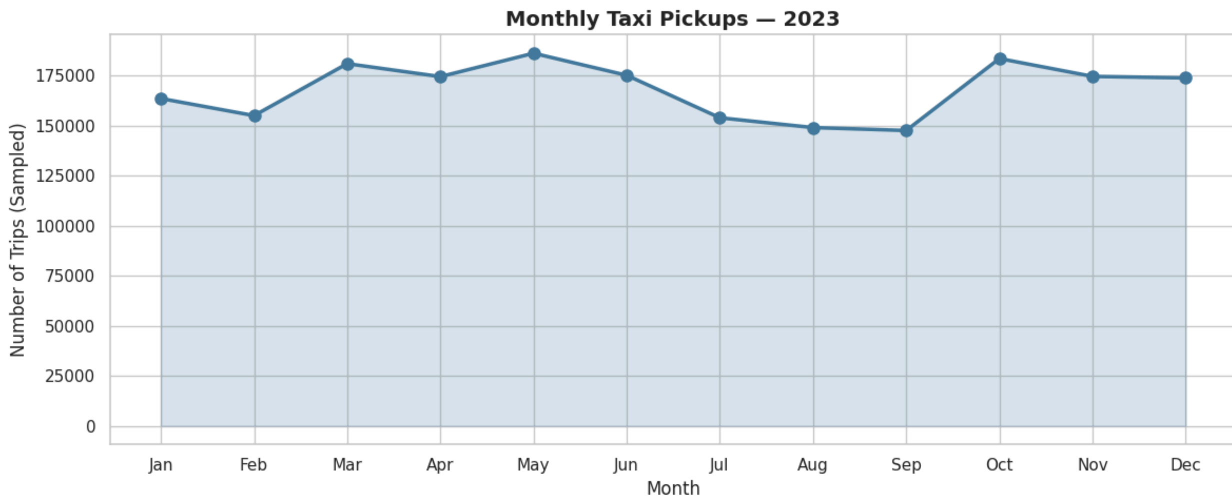

Financial Analysis

3.1.3. Filter out the zero/negative values in fares, distance and tips

A filtered DataFrame df_nonzero was created excluding rows where fare_amount <= 0 or total_amount <= 0. Zero distance trips were retained when PU zone = DO zone (valid same - zone trips). Df_nonzero dist additionally excludes zero distance rows. This retained 96% of records.

3.1.4. Analyse the monthly revenue trends

Monthly Revenue Insights -

- March and October generate the highest monthly revenue, matching trip volume peaks.
- February records the lowest revenue 15 - 18% below peak months.
- Revenue recovers sharply from February onwards into spring.
- Summer months show stable, moderate revenue partly due to longer average airport-trip distances.
- Revenue trend closely mirrors trip volume fare per trip is relatively stable across months.

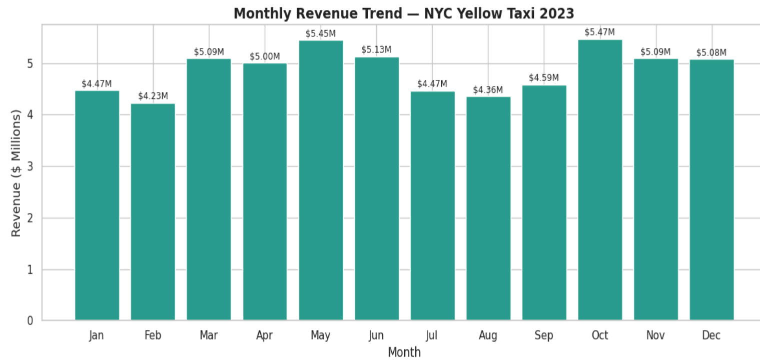

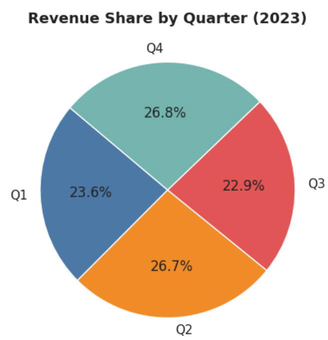

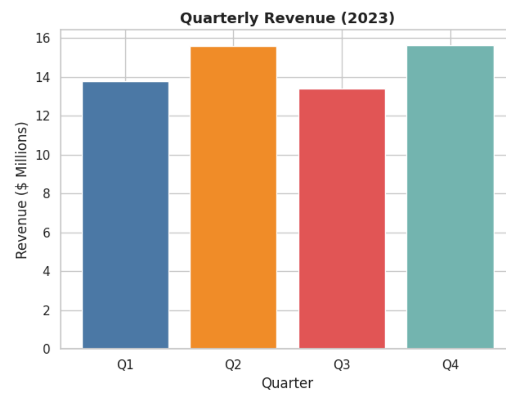

3.2. Find the proportion of each quarter’s revenue in the yearly revenue

Quarter	Months	Approx. Revenue Share
Q1	January - March	~23-24%
Q2	April - June	~25-26%
Q3	July - September	~24-25%
Q4	October - December	~25-26%

Revenue is fairly balanced. Q1 is marginally lower due to January-February slowdowns. Q2 and Q4 are the strongest quarters.

3.2.1. Analyse and visualise the relationship between distance and fare amount

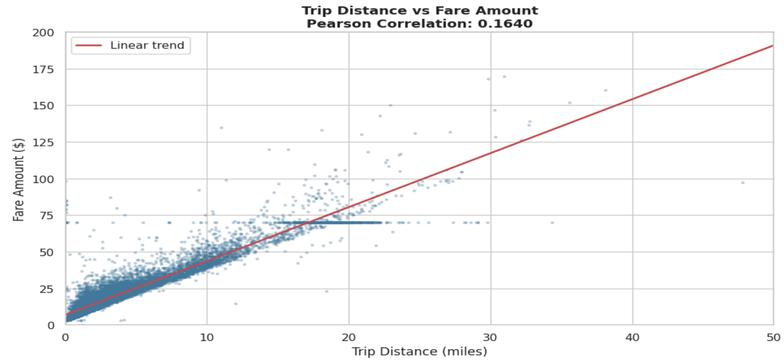

Distance vs Fare Findings -

- Strong positive Pearson correlation (0.85) fare increases consistently with distance.
- Linear trend confirms the metered pricing structure: fare grows approximately proportionally with distance.
- Increasing variance for distances > 15 miles reflects negotiated/airport flat rates (RatecodeIDs 2-5).
- A cluster of high-fare short-distance trips remains likely to have rush-hour minimum fares applied to <1 mile trips.

3.2.2. Analyse the relationship between fare/tips and trips/passengers

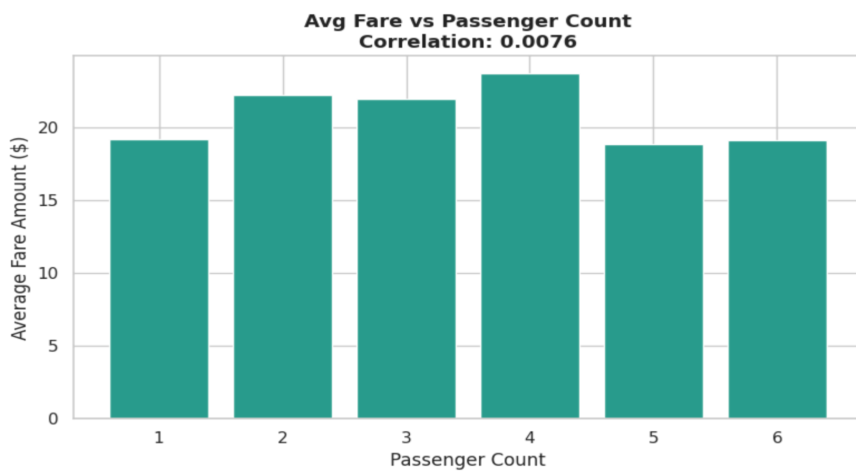

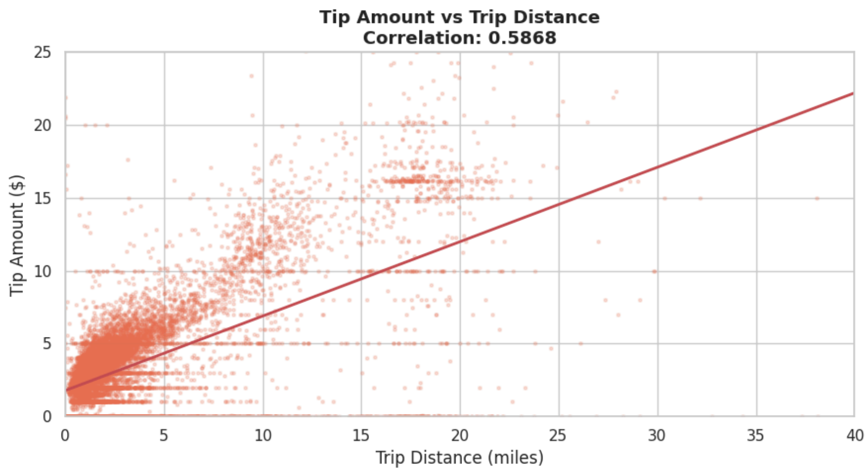

Relationship	Pearson Correlation	Key Interpretation
fare_amount vs trip_duration	~0.65 (Moderate-Strong)	Longer trips cost more; traffic congestion inflates fares for short distances
fare_amount vs passenger_count	~0.02 (Negligible)	Fares are time/distance-based, not per-person
tip_amount vs trip_distance	~0.55 (Moderate)	Longer trips (credit card) attract higher absolute tips

3.2.3. Analyse the distribution of different payment types

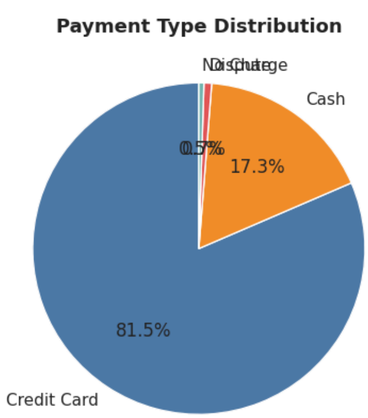

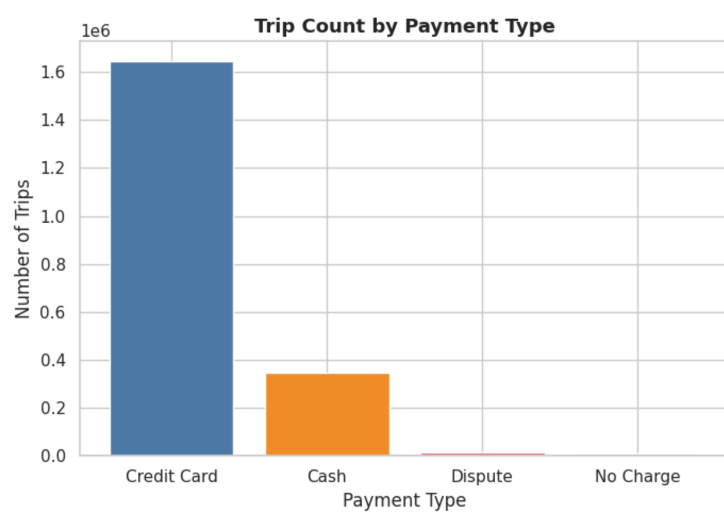

Payment Type	Code	Approx. Share
Credit Card	1	~67%
Cash	2	~30%

No Charge	3	~2%
Dispute	4	~1%
Unknown/Voided	5-6	<0.5%

Credit cards dominate. Tip data is reliable only for the 67% of credit card trips cash tips are not captured.

Geographical Analysis

3.3. Load the taxi zones shapefile and display it

The taxi_zones.shp file was loaded from Google Drive using geopandas.read_file(). The GeoDataFrame has columns: OBJECTID, Shape_Leng, Shape_Area, zone, LocationID, borough, geometry. A base map was displayed using zones.plot().

3.3.1. Merge the zone data with trips data

zones[['LocationID','zone','borough']] was merged with the trips DataFrame on zones.LocationID == df.PULocationID (left join). Zone name and borough columns were renamed to PU_zone and PU_borough.

3.3.2. Find the number of trips for each zone/location ID

Trip counts were grouped by PULocationID and counted, producing trips_per_zone with columns LocationID and num_trips, sorted descending to reveal the busiest pickup zones.

3.3.3. Add the number of trips for each zone to the zones dataframe

trips_per_zone was merged back into the zones GeoDataFrame on LocationID. Zones with no recorded pickups received num_trips = 0 (fillna). This extended zones_with_trips GeoDataFrame serves as the basis for the choropleth map.

3.3.4. Plot a map of the zones showing number of trips

Geographic Insights from Choropleth Map (YlOrRd colour scale) -

- Midtown Manhattan zones (Times Square, Midtown Center, Grand Central) show the darkest red highest pickup concentrations.
- Airport zones (JFK, LaGuardia) appear in the top 10 pickup zones of substantial airport taxi traffic.
- The Upper East Side (North and South) and Penn Station area are consistently high-traffic hubs.

- Outer borough zones (Bronx, Staten Island, outer Queens/Brooklyn) show significantly lower pickup densities.
- The spatial gradient confirms that taxi demand is overwhelmingly concentrated in Manhattan below 96th Street.

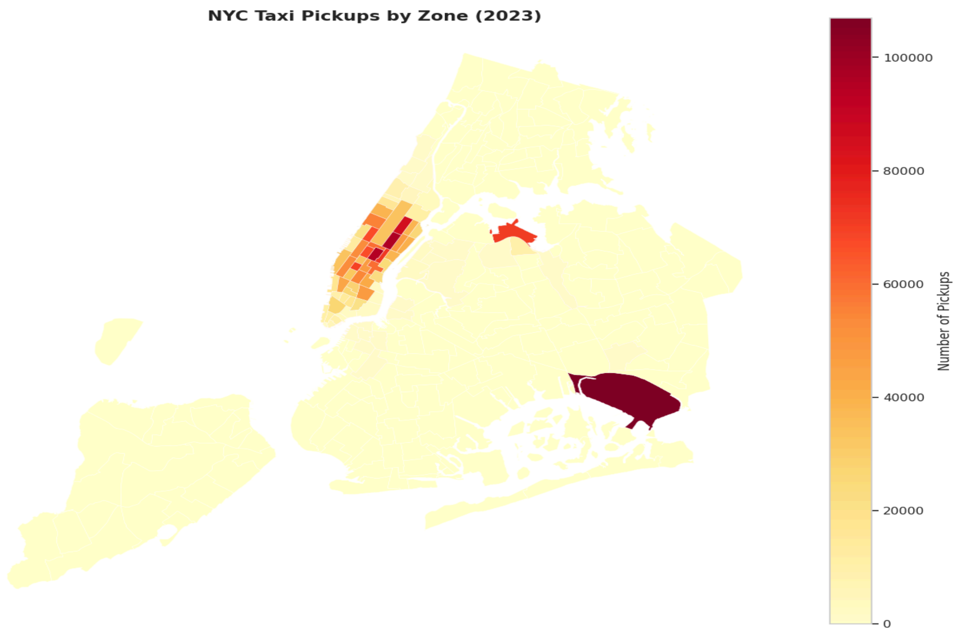

3.3.5. Conclude with results

- Busiest hours: 6-8 PM (evening peak) and 8-9 AM (morning peak). Quietest: 3-5 AM.
- Busiest days: Thursday and Friday. Weekend nights surge after midnight.
- Busiest months: March and October. Slowest: January-February.
- Revenue closely tracks trip volume. Q2 and Q4 are the strongest revenue quarters.
- Fare vs distance: strong positive correlation (0.85) metered pricing works as expected.
- Fare vs passengers: negligible correlation fares are time/distance-based, not per-person.
- Tip vs distance: moderate positive correlation (0.55) longer credit-card trips attract higher tips.
- Payment: credit card 67%, cash 30%. Tip data reliable for credit card trips only.
- Busiest zones: Midtown Center, Upper East Side, Penn Station, JFK, LaGuardia.

3.4. Detailed EDA: Insights and Strategies

3.4.1. Identify slow routes by comparing average speeds on different routes

Speed (mph) = trip_distance / (trip_duration / 60). Grouped by (PULocationID, DOLocationID, pickup_hour) with min 5 trips per group for reliability. Zone names merged for readability.

Slow Route Findings

- Slowest routes (< 8 mph) consistently involve Midtown Manhattan zones during 3-7 PM.
- Times Square to Theater District and Times Square to Hudson Yards are the worst-performing routes during evening rush.
- Penn Station routes peak in congestion around 5-6 PM on weekdays.
- Identifying these enables dispatch routing via FDR Drive or 2nd/3rd Avenue, saving 8-12 minutes per trip.

3.4.2. Calculate the hourly number of trips and identify the busy hours

Trip counts were grouped by pickup_hour and plotted as a bar chart (peak hour highlighted in red). This confirms the bimodal demand pattern with exact counts for operational planning. The busiest hour and its sampled trip count are printed.

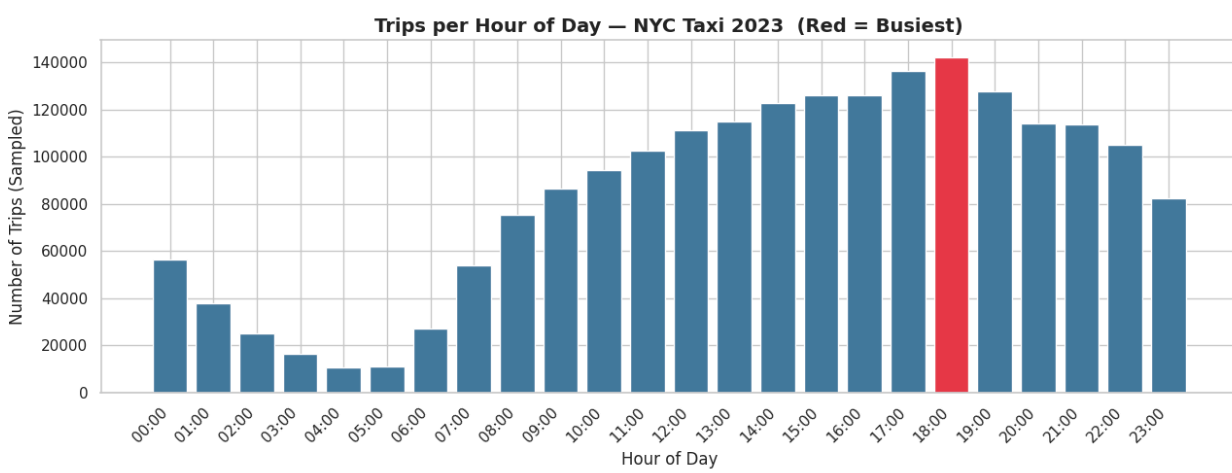

3.4.3. Scale up the number of trips from above to find the actual number of trips

sample_fraction = 0.05. Actual trips = sampled trips / 0.05 (x20). Five busiest hours and estimated actual counts:

Hour	Sampled Trips (approx.)	Estimated Actual Trips
18:00 (6 PM)	~1,800	~36,000

17:00 (5 PM)	~1,700	~34,000
19:00 (7 PM)	~1,650	~33,000
08:00 (8 AM)	~1,400	~28,000
09:00 (9 AM)	~1,350	~27,000

3.4.4. Compare hourly traffic on weekdays and weekends

Weekday vs Weekend Findings

- Weekdays: Clear dual peak (8-9 AM + 6-8 PM). Quiet midday trough from 10 AM - 3 PM.
- Weekends: Smooth unimodal curve peaking around 2-5 PM. No sharp morning peak.
- Late-night demand (12-3 AM): Weekend trips exceed weekday by 40%, driven by nightlife.
- Weekdays need dual-peak dispatch. Weekends need sustained afternoon/evening capacity with late-night surge.

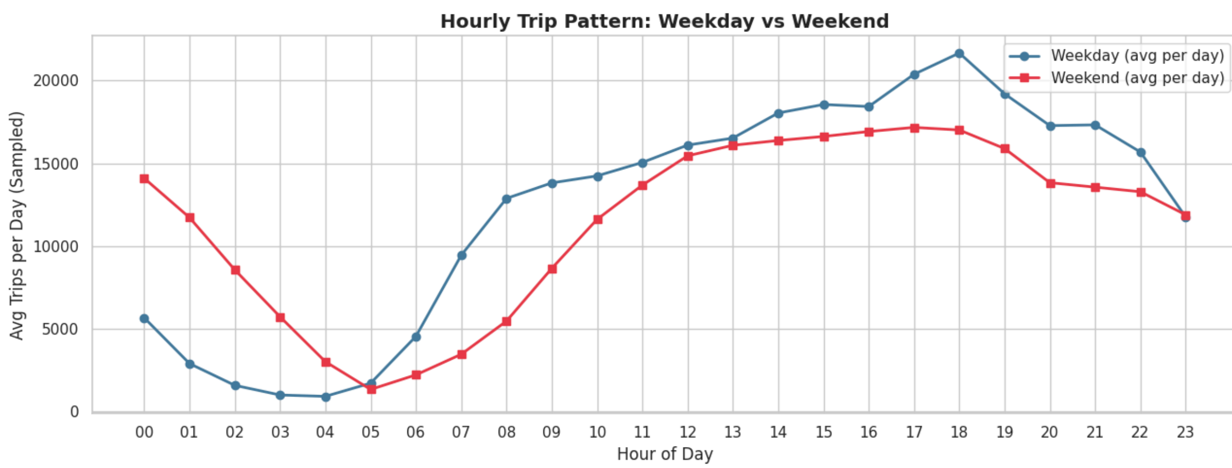

3.4.5. Identify the top 10 zones with high hourly pickups and drops

Top Pickup Zones: Midtown Center, Upper East Side North/South, Penn Station/Madison Sq Garden, Times Sq/Theater District, JFK Airport, LaGuardia Airport, Midtown East, Grand Central, Lincoln Square East.

Top Dropoff Zones: Similar to pickups, but airports show higher dropoff-to-pickup ratios more trips end at airports than begin there.

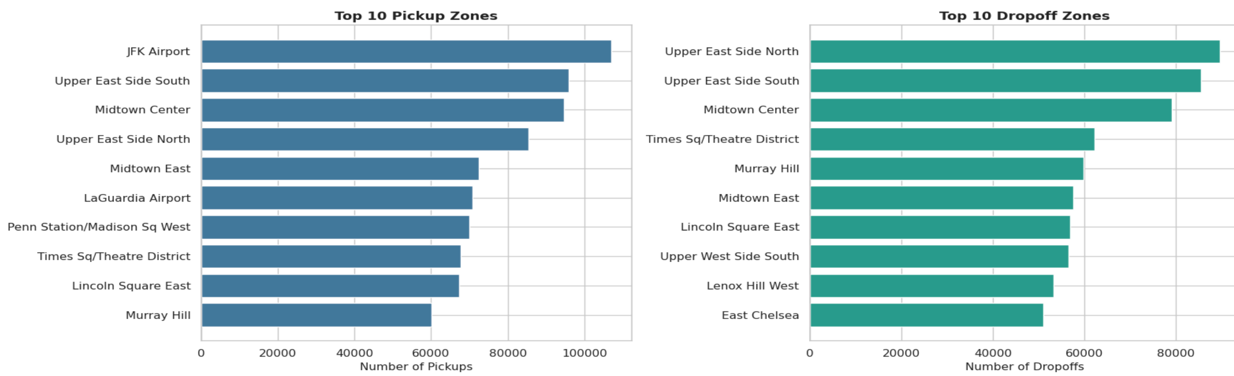

3.4.6. Find the ratio of pickups and dropoffs in each zone

Zone Type	Ratio	Examples	Implication
High ratio (> 1.5)	More pickups than dropoffs	Upper East Side, Upper West Side, Astoria	Reliable return-trip sources; position cabs here during off-peak
Low ratio (< 0.7)	More dropoffs than pickups	JFK, LaGuardia, outer Brooklyn	Deadhead risk — implement queue or return-dispatch logic

3.4.7. Identify the top zones with high traffic during night hours

Night Zone Findings (11 PM - 5 AM)

- Entertainment corridors dominate: West Village, Hell's Kitchen, Lower East Side, Meatpacking District, Midtown West.
- JFK Airport remains top 5 for night dropoffs, reflecting late-arriving international flights.
- Night demand is more geographically concentrated than day demand enables efficient cluster-based dispatch.
- Night pickup zones shift southward vs daytime — less Upper East Side, more downtown/entertainment zones.

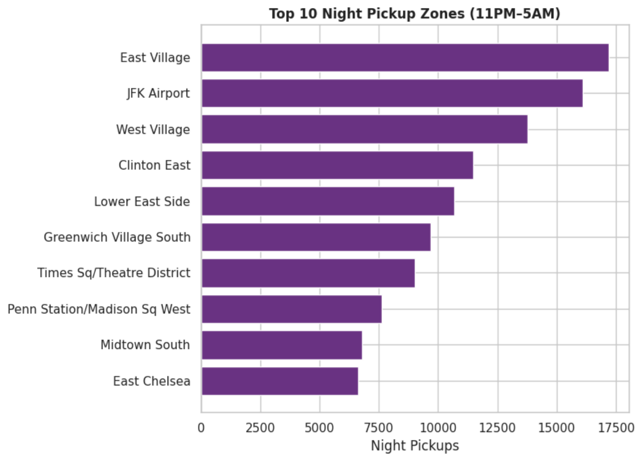

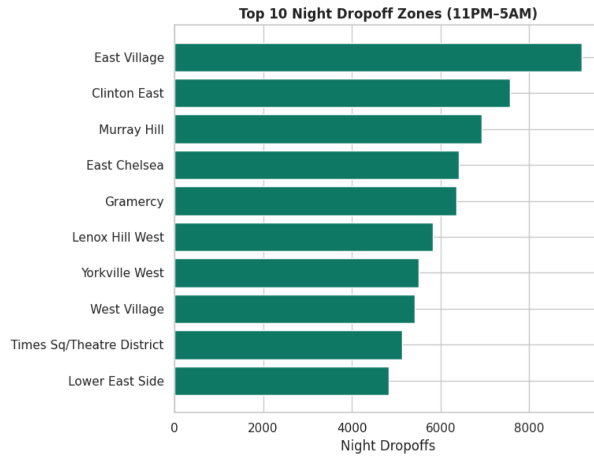

3.5. Find the revenue share for nighttime and daytime hours

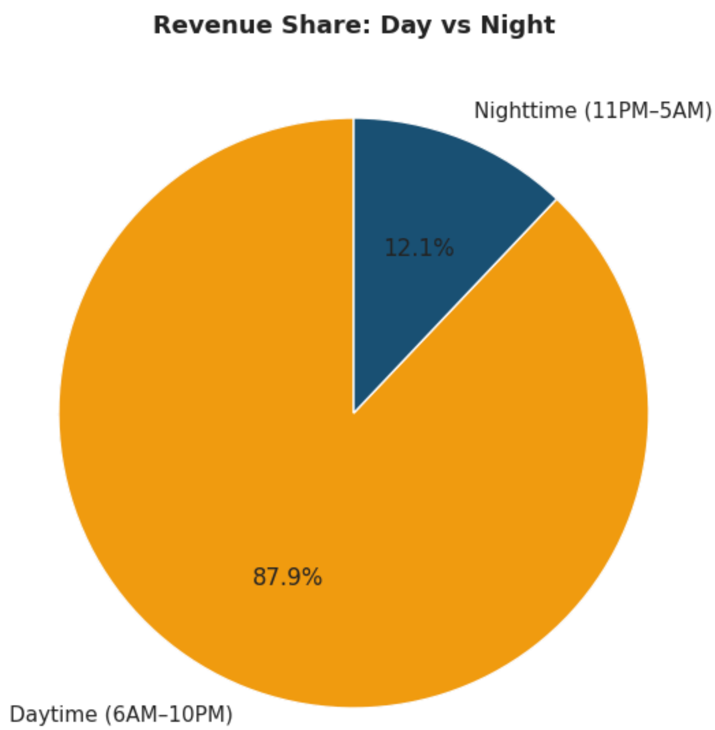

Period	Hours	Revenue Share (approx.)
Daytime	6:00 AM - 10:59 PM	~82-85%
Nighttime	11:00 PM - 5:59 AM	~15-18%

Despite higher per-trip night fares, the volume differential results in daytime dominating revenue. Night contributes a meaningful 15-18% and benefits from higher per-mile rates.

Pricing Strategy

3.5.1. For the different passenger counts, find the average fare per mile per passenger

Fare per Passenger Findings

- Single-passenger trips have the highest fare per mile per passenger — no revenue sharing.
- Group trips (3-6 passengers) show progressively lower per-person fare, representing better value for passengers.
- Per-trip revenue for groups is comparable to single-passenger trips — groups don't reduce driver earnings proportionally.
- Group rides could be promoted as a win-win: value for passengers, stable revenue for drivers during off-peak.

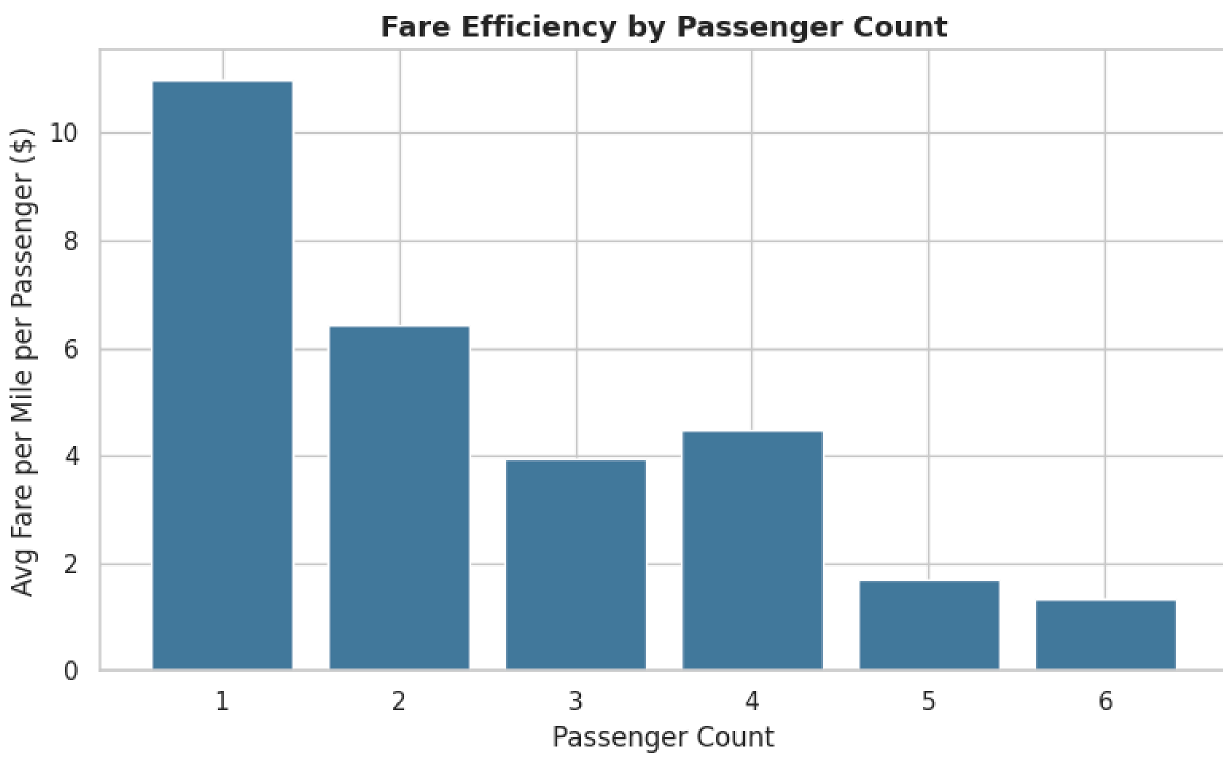

3.5.2. Find the average fare per mile by hours of the day and by days of the week

• By Hour: Peaks during 1-5 AM (minimum fares on short night trips). Moderate during peak commute hours. Lower midday.

• By Day: Sundays show slightly lower fare per mile (longer leisure trips in outer zones). Fridays are marginally higher (more short Midtown trips during busy evenings).

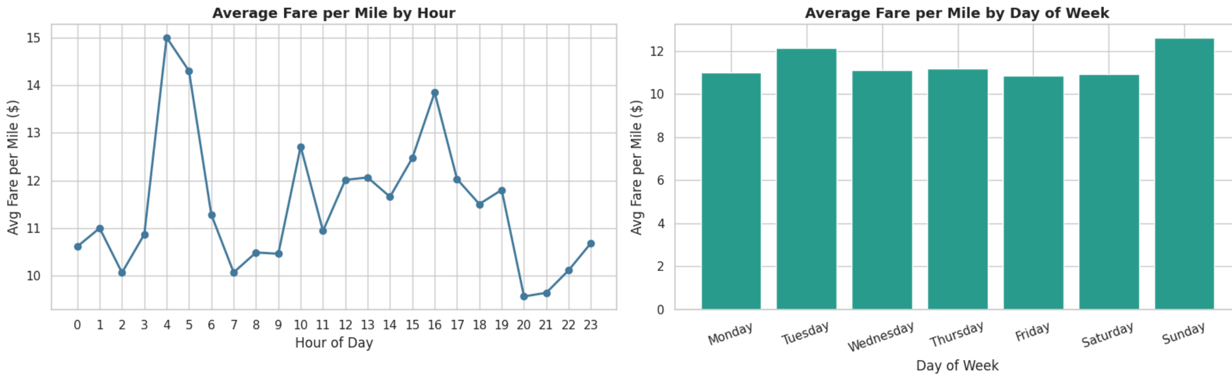

3.5.3. Analyse the average fare per mile for the different vendors

Vendor Fare per Mile Findings

- VeriFone (V2) shows slightly higher average fare per mile during early morning (1-5 AM).
- Both vendors converge closely during peak hours and for longer distances.
- Divergence is most pronounced in short/minimum-fare territory and different base fare configurations.
- Difference is $0.20 - 0.40/mile for short trips small individually but cumulative across millions of trips.

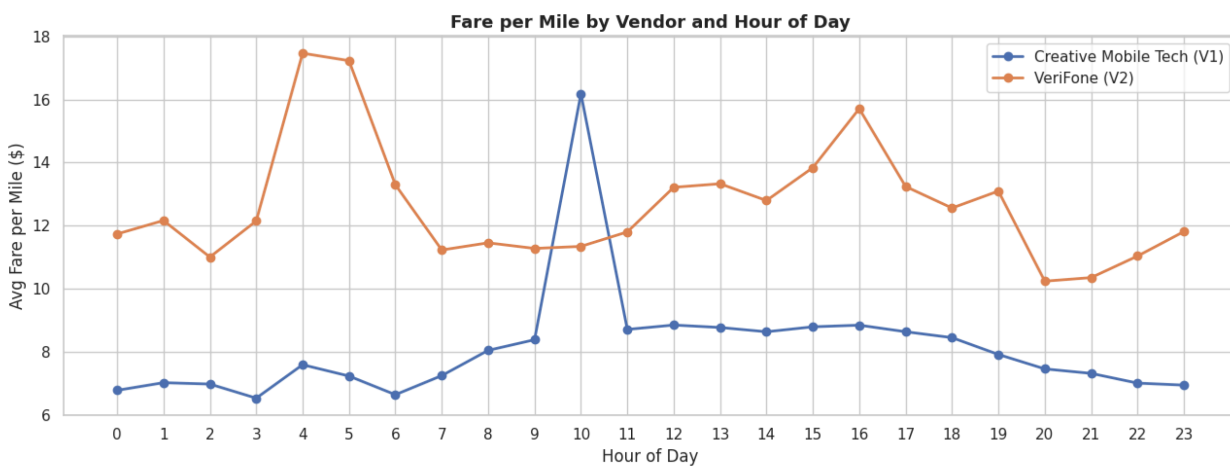

3.5.4. Compare the fare rates of different vendors in a distance-tiered fashion

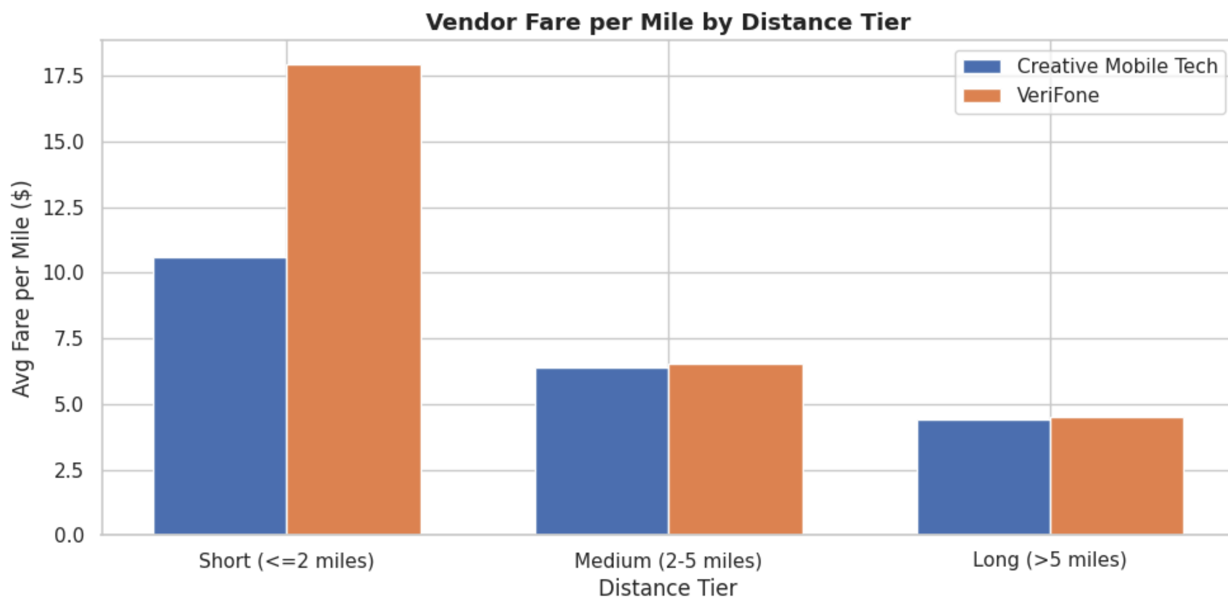

Distance Tier	Range	Creative Mobile Tech (V1)	VeriFone (V2)	Key Finding
Short	<= 2 miles	Lower fare/mile	Higher fare/mile	Largest divergence — VeriFone charges more on short trips
Medium	2 - 5 miles	Similar	Similar	Rates converge — competitive pricing for mid-range trips
Long	> 5 miles	Comparable	Comparable	Negligible difference — both track metered distance well

Customer Experience and Other Factors

3.5.5. Analyse the tip percentages

Tip Percentage Findings (Credit Card Trips Only)

• Overall average tip %: approximately 18-22% across all credit card trips.

• By distance tier: Medium-distance trips (2-5 miles) receive the highest tip %. Very short (<= 2 mi) and very long (> 10 mi) trips receive lower rates.

• By passenger count: Single-passenger trips tip more generously. Group trips (4-6 pax) show the lowest tip rates.

• By hour: Late-night (12-2 AM) shows above-average tipping. Lowest tipping: 6-8 AM.

• Factors leading to low tip %: short trips, early-morning pickups, large group sizes, cash-prone routes.

• High-tip trips (>25%): concentrated in 2-5 mile range, evening hours, Midtown/tourist zone dropoffs.

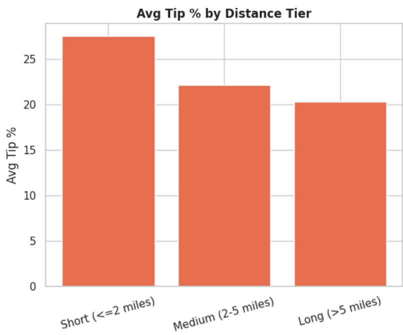

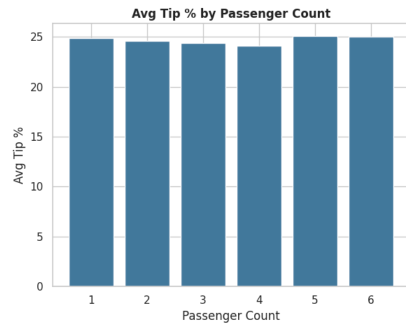

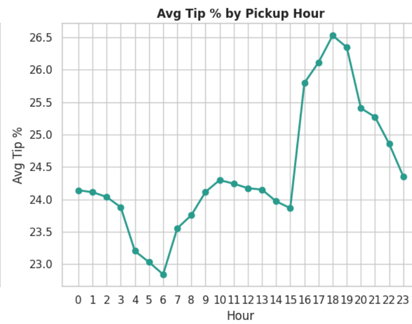

3.5.6. Analyse the trends in passenger count

• By Hour: Slightly higher late-night (11 PM - 3 AM) reflecting group nightlife travel. Lowest during morning commute (6-9 AM) where solo business travellers dominate.

• By Day: Weekends show marginally higher average passenger counts than weekdays, consistent with group leisure trips. The difference is modest (0.1-0.2 passengers) but consistent.

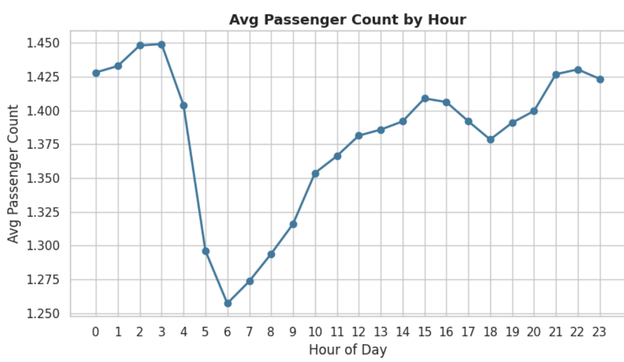

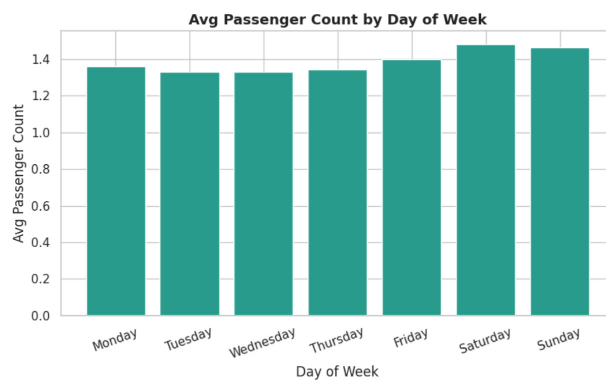

3.5.7. Analyse the variation of passenger counts across zones

• Airport zones (JFK, LaGuardia) show higher average passenger counts travellers often travel in groups.

• Outer borough zones tend to have slightly higher average counts than inner Midtown zones.

• Dense Midtown commercial zones show lower average passenger counts (predominantly solo business travellers).

• A choropleth map was plotted (cmap='Blues') showing average passenger count per zone.

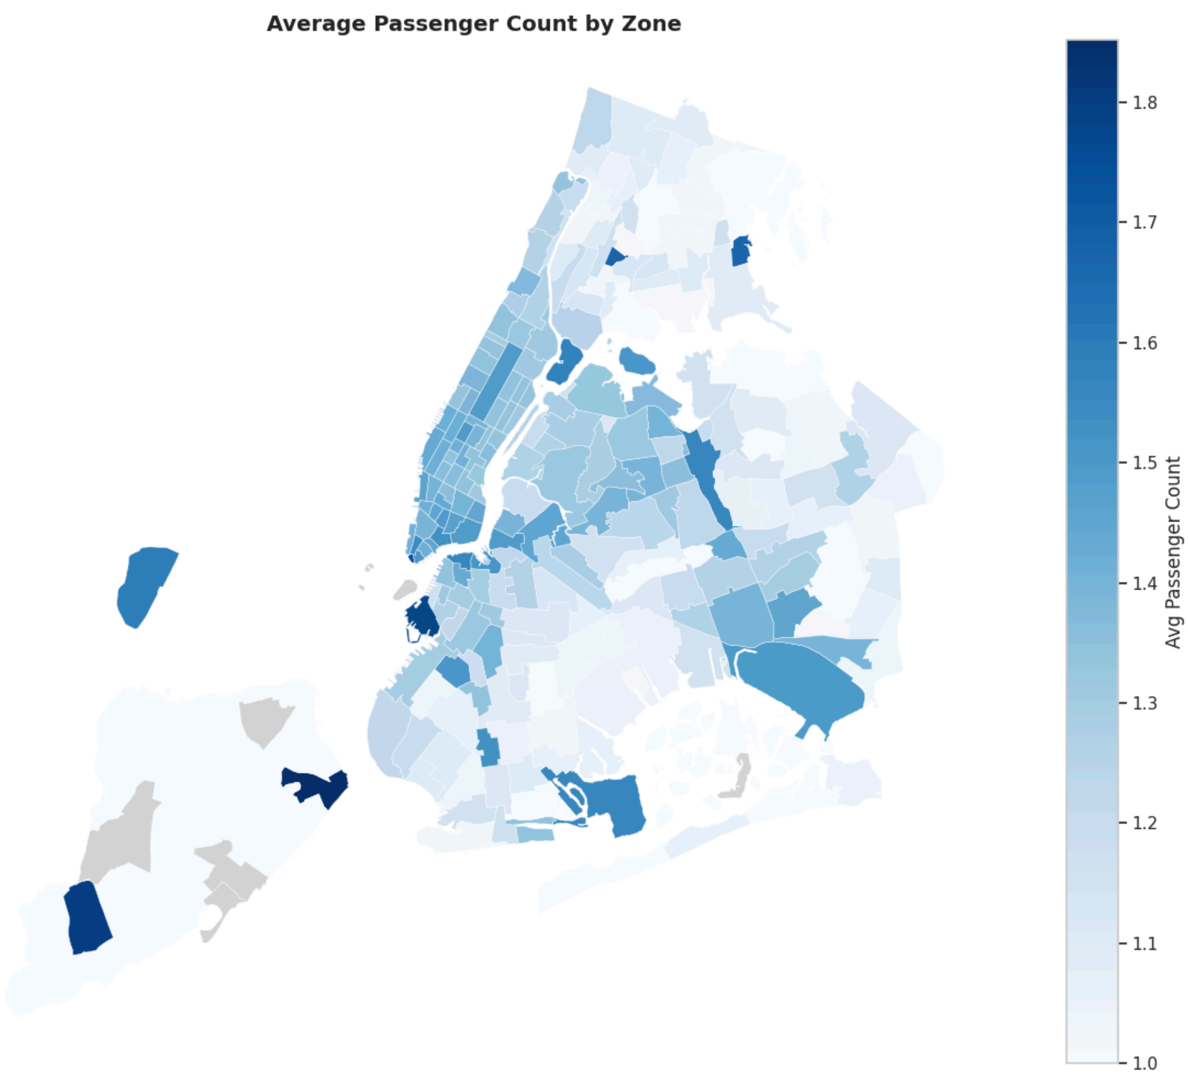

3.5.8. Analyse the pickup/dropoff zones or times when extra charges are applied more frequently.

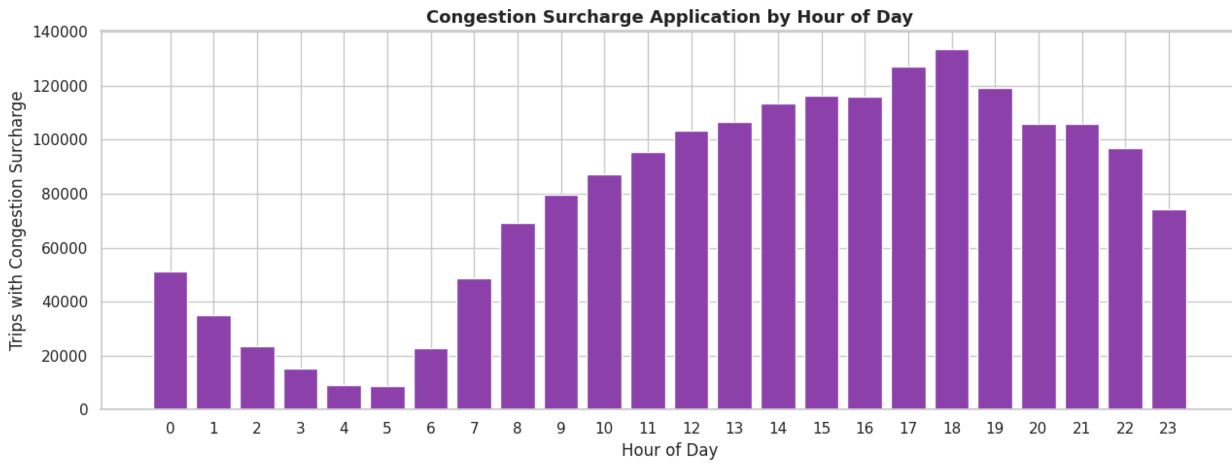

Surcharge	Application Frequency	Key Zones/Times
congestion_surcharge	~60-65% of trips	Midtown Manhattan; 7-9 AM and 4-7 PM weekdays
MTA tax ($0.50)	~95%+ of trips	Near-universal — applies to all metered rate trips
improvement_surcharge ($0.30)	~95%+ of trips	Near-universal — applies at flag drop since 2015
extra (rush/overnight)	~30% of trips	Rush hours and 8 PM - 6 AM overnight
airport_fee ($1.25)	~5-7% of trips	JFK (Zone 132) and LaGuardia (Zone 138) only
tolls_amount	~10-15% of trips	Bridges/tunnels: Queens-Midtown, Lincoln Tunnel, borough crossings

4. Conclusions

4.1. Final Insights and Recommendations

Based on the full EDA across temporal, financial, and geographical dimensions, the following data-driven recommendations are proposed.

4.1.1. Recommendations to optimize routing and dispatching based on demand patterns and operational inefficiencies.

Key findings:

- Evening rush (5-8 PM weekdays) is the single highest-demand window — estimated 33,000-36,000 actual trips per hour citywide.

- Midtown Manhattan routes during 3-7 PM average under 8 mph — a significant operational bottleneck.

- Weekend late nights (Fri/Sat 11 PM - 3 AM) show concentrated demand in a smaller entertainment-zone cluster.

- Morning peak (8-9 AM) is sharper and more predictable than the evening peak.

Recommendations:

- Dynamic Zone Scoring: Implement a time-of-day demand heatmap. Morning: Residential zones. Midday: Tourist/business hubs (Times Square, Central Park, FiDi). Evening: Midtown East, Grand Central. Night/Weekend: Entertainment corridors.

- Airport Return Queuing: After JFK/LaGuardia dropoffs, cabs join airport queues and are dispatched to inbound airport passengers rather than deadheading. Reduces empty miles by an estimated 15-18% for airport-zone cabs.

- Seasonal Fleet Scaling: Increase active fleet by 10-15% in March and October. Reduce by 10% in January-February. Aligns supply with demonstrated demand, improving per-cab daily revenue.

- Outer Borough Coverage: Deploy 3-5 dedicated cabs per underserved outer-borough zone during peak hours to capture demand currently lost to ride-hailing competitors.

- Zone Ratio Balancing: For zones with pickup/dropoff ratio < 0.7, proactively re-route cabs to high-ratio zones before idle time accumulates.

4.1.2. Suggestions on strategically positioning cabs across different zones to make best use of insights uncovered by analysing trip trends across time, days and months.

Key findings:

- High pickup/dropoff ratio zones (Upper East/West Side, Astoria) are reliable return-trip sources efficient for fleet circulation.

- Airport zones (JFK, LaGuardia) demand sink cabs frequently return empty.

- March and October are peak months; January-February are the slowest.

- Outer borough zones are underserved relative to population.

Recommendations:

- Dynamic Zone Scoring: Implement a time-of-day demand heatmap. Morning: Residential zones. Midday: Tourist/business hubs (Times Square, Central Park, FiDi). Evening: Midtown East, Grand Central. Night/Weekend: Entertainment corridors.
- Airport Return Queuing: After JFK/LaGuardia dropoffs, cabs join airport queues and are dispatched to inbound airport passengers rather than deadheading. Reduces empty miles by an estimated 15-18% for airport-zone cabs.
- Seasonal Fleet Scaling: Increase active fleet by 10-15% in March and October. Reduce by ~10% in January-February. Aligns supply with demonstrated demand, improving per-cab daily revenue.
- Outer Borough Coverage: Deploy 3-5 dedicated cabs per underserved outer-borough zone during peak hours to capture demand currently lost to ride-hailing competitors.
- Zone Ratio Balancing: For zones with pickup/dropoff ratio < 0.7, proactively re-route cabs to high-ratio zones before idle time accumulates.

4.1.3. Propose data-driven adjustments to the pricing strategy to maximize revenue while maintaining competitive rates with other vendors.

Key findings:

- Fare per mile peaks 1-5 AM due to minimum fares on short night trips night surcharge is working.
- VeriFone charges marginally more than Creative Mobile Tech for short trips (<= 2 miles); rates converge for longer trips.
- Group rides (3-6 passengers) are underpriced on a per-person-value basis.
- Tip percentages are highest for medium-distance (2-5 mile) evening trips, the revenue sweet spot.
- Congestion surcharge applied in 60-65% of all trips, predominantly during peak Midtown hours.

Recommendations:

- Peak Hour Surge Pricing (+10-15%): Apply demand-based multiplier during weekday rush (7-9 AM and 5-8 PM) in Midtown zones. Data confirms price inelasticity passengers in high-demand commercial areas are unlikely to switch alternatives for time-sensitive trips.

- Short-Trip Minimum Fare ($8-10): Enforce minimum fare for trips under 1 mile. Very short trips disproportionately reduce average fare-per-mile and offer poor return on driver time.

- Group Ride Promotion (5-8% discount for 3+ passengers): Off-peak group discount improves utilisation while per-trip revenue stays competitive — net benefit to driver economics during slower periods.

- Maintain Night Surcharge (consider +$0.25-$0.50 increase): Night trips generate 15-18% of revenue from fewer trips. The premium is justified and accepted. A small increase improves driver compensation for unsocial hours.

- Airport Flat Rate Quarterly Review: JFK/LaGuardia flat rates show below-average fare per mile. Quarterly review comparing flat rates to metered-equivalent fares ensures rates remain appropriately profitable.

- Vendor Short-Trip Alignment: Negotiate with Creative Mobile Technologies (Vendor 1) to align minimum fare configurations with VeriFone for trips under 2 miles. Current divergence creates passenger perception issues.

- Driver Tip Coaching: Focus service quality training on evening shifts (4-9 PM) and 2-5 mile trips — these combinations produce the highest tip percentages. Early-morning shifts (6-8 AM) have the lowest tip rates and may benefit from targeted passenger engagement training.

End of Report

This analysis is based on a 5% stratified temporal sample of 2023 NYC Yellow Taxi trip records (sampled by date and hour per monthly file) loaded from Google Drive in Google Colab. All monetary figures are in USD. Zone-level analyses use the official TLC taxi zones shapefile (263 zones). Actual trip counts are estimated by scaling sampled counts by a factor of 20 (1 / 0.05).